

Intermediate Quantitative Methods

Regression anatomy

Kevin Munger

Week 2, part 2

European University Institute, SPS

Today

- What least squares does

Today

- What least squares does
- What “holding constant” means

Today

- What least squares does
- What “holding constant” means
- Uncertainty in a slope

Today

- What least squares does
- What “holding constant” means
- Uncertainty in a slope
- Assumptions, and what happens when they break

Today

- What least squares does
- What “holding constant” means
- Uncertainty in a slope
- Assumptions, and what happens when they break
- Interactions

What least squares does

The pieces

The pieces

- Model: $y_i = \alpha + \beta x_i + \varepsilon_i$

The pieces

- Model: $y_i = \alpha + \beta x_i + \varepsilon_i$
- α : intercept, the expected y when $x = 0$

The pieces

- Model: $y_i = \alpha + \beta x_i + \varepsilon_i$
- α : intercept, the expected y when $x = 0$
- β : slope, the difference in expected y per unit of x


The pieces

- Model: $y_i = \alpha + \beta x_i + \varepsilon_i$
- α : intercept, the expected y when $x = 0$
- β : slope, the difference in expected y per unit of x
- Fitted value: $\hat{y}_i = \hat{\alpha} + \hat{\beta} x_i$, the point on the line

The pieces

- Model: $y_i = \alpha + \beta x_i + \varepsilon_i$
- α : intercept, the expected y when $x = 0$
- β : slope, the difference in expected y per unit of x
- Fitted value: $\hat{y}_i = \hat{\alpha} + \hat{\beta}x_i$, the point on the line
- Residual: $e_i = y_i - \hat{y}_i$, the vertical distance to the line

A fitted line and its residuals, $n = 40$



Ordinary least squares (OLS)

Ordinary least squares (OLS)

- Choose $\hat{\alpha}$ and $\hat{\beta}$ to minimize $\sum_i e_i^2$

Ordinary least squares (OLS)

- Choose $\hat{\alpha}$ and $\hat{\beta}$ to minimize $\sum_i e_i^2$
- Solution: $\hat{\beta} = \text{Cov}(x, y) / \text{Var}(x)$

Ordinary least squares (OLS)

- Choose $\hat{\alpha}$ and $\hat{\beta}$ to minimize $\sum_i e_i^2$
- Solution: $\hat{\beta} = \text{Cov}(x, y) / \text{Var}(x)$
- The line passes through $(\bar{x}, \bar{y})$


Ordinary least squares (OLS)

- Choose $\hat{\alpha}$ and $\hat{\beta}$ to minimize $\sum_i e_i^2$
- Solution: $\hat{\beta} = \text{Cov}(x, y) / \text{Var}(x)$
- The line passes through $(\bar{x}, \bar{y})$
- Residuals average zero and are uncorrelated with x , by construction

Ordinary least squares (OLS)

- Choose $\hat{\alpha}$ and $\hat{\beta}$ to minimize $\sum_i e_i^2$
- Solution: $\hat{\beta} = \text{Cov}(x, y) / \text{Var}(x)$
- The line passes through $(\bar{x}, \bar{y})$
- Residuals average zero and are uncorrelated with x , by construction
- So residuals cannot tell you whether x is confounded

A binary predictor, 100 per group



Reading coefficients

- Binary x : $\hat{\beta}$ is the difference in group means, $\hat{\alpha}$ the mean of the $x = 0$ group

Reading coefficients

- Binary x : $\hat{\beta}$ is the difference in group means, $\hat{\alpha}$ the mean of the $x = 0$ group
- Continuous x : difference in average y per unit of x

Reading coefficients

- Binary x : $\hat{\beta}$ is the difference in group means, $\hat{\alpha}$ the mean of the $x = 0$ group
- Continuous x : difference in average y per unit of x
- Units matter: years or decades, euros or thousands of euros

Reading coefficients

- Binary x : $\hat{\beta}$ is the difference in group means, $\hat{\alpha}$ the mean of the $x = 0$ group
- Continuous x : difference in average y per unit of x
- Units matter: years or decades, euros or thousands of euros
- Intercept at $x = 0$ often describes nobody: center x at its mean

Reading coefficients

- Binary x : $\hat{\beta}$ is the difference in group means, $\hat{\alpha}$ the mean of the $x = 0$ group
- Continuous x : difference in average y per unit of x
- Units matter: years or decades, euros or thousands of euros
- Intercept at $x = 0$ often describes nobody: center x at its mean
- Monday's lab does all of this in R

Holding constant

Multiple regression

Multiple regression

- $y_i = \alpha + \beta_1 x_{1i} + \beta_2 x_{2i} + \varepsilon_i$

Multiple regression

- $y_i = \alpha + \beta_1 x_{1i} + \beta_2 x_{2i} + \varepsilon_i$
- β_1 : compares units that differ in x_1 and have the same x_2

Multiple regression

- $y_i = \alpha + \beta_1 x_{1i} + \beta_2 x_{2i} + \varepsilon_i$
- β_1 : compares units that differ in x_1 and have the same x_2
- “Holding constant”, “controlling for”, “adjusting for”: the same thing

Multiple regression

- $y_i = \alpha + \beta_1 x_{1i} + \beta_2 x_{2i} + \varepsilon_i$
- β_1 : compares units that differ in x_1 and have the same x_2
- “Holding constant”, “controlling for”, “adjusting for”: the same thing
- Nothing is physically held constant

- Step 1: regress x_1 on x_2 , keep the residuals $\tilde{x}_1$

Regression anatomy

- Step 1: regress x_1 on x_2 , keep the residuals $\tilde{x}_1$
- $\tilde{x}_1$: the part of x_1 that x_2 does not predict


Regression anatomy

- Step 1: regress x_1 on x_2 , keep the residuals $\tilde{x}_1$
- $\tilde{x}_1$: the part of x_1 that x_2 does not predict
- Step 2: regress y on $\tilde{x}_1$

- Step 1: regress x_1 on x_2 , keep the residuals $\tilde{x}_1$
- $\tilde{x}_1$: the part of x_1 that x_2 does not predict
- Step 2: regress y on $\tilde{x}_1$
- The slope equals $\hat{\beta}_1$ from the full regression, exactly

- Step 1: regress x_1 on x_2 , keep the residuals $\tilde{x}_1$
- $\tilde{x}_1$: the part of x_1 that x_2 does not predict
- Step 2: regress y on $\tilde{x}_1$
- The slope equals $\hat{\beta}_1$ from the full regression, exactly
- Frisch–Waugh–Lovell theorem

Simulated data, true $\beta_1 = 0.5$, $n = 1,500$



What follows from this

What follows from this

- $\hat{\beta}_1$ uses only the leftover variation in x_1

What follows from this

- $\hat{\beta}_1$ uses only the leftover variation in x_1
- Here: slope 1.34 without the control, 0.56 with it

What follows from this

- $\hat{\beta}_1$ uses only the leftover variation in x_1
- Here: slope 1.34 without the control, 0.56 with it
- Highly correlated predictors leave little variation: large standard errors

What follows from this

- $\hat{\beta}_1$ uses only the leftover variation in x_1
- Here: slope 1.34 without the control, 0.56 with it
- Highly correlated predictors leave little variation: large standard errors
- That is multicollinearity: a data limitation, not a violation

What follows from this

- $\hat{\beta}_1$ uses only the leftover variation in x_1
- Here: slope 1.34 without the control, 0.56 with it
- Highly correlated predictors leave little variation: large standard errors
- That is multicollinearity: a data limitation, not a violation
- Every added control narrows the comparison being made

Uncertainty in a slope

$\hat{\beta}$ is an estimate

$\hat{\beta}$ is an estimate

- A different sample gives a different line


$\hat{\beta}$ is an estimate

- A different sample gives a different line
- $\hat{\beta}$ has a sampling distribution, like $\bar{x}$ last week


$\hat{\beta}$ is an estimate

- A different sample gives a different line
- $\hat{\beta}$ has a sampling distribution, like $\bar{x}$ last week
- Simulation: true slope 0.5, $n = 100$, 1,000 datasets

Fitted lines from 100 simulated datasets



Slopes from 1,000 simulated datasets



Standard error of a slope

Standard error of a slope

- $SE(\hat{\beta}) \approx \frac{\sigma_{\varepsilon}}{\sqrt{n} s_x}$

Standard error of a slope

- $SE(\hat{\beta}) \approx \frac{\sigma_{\varepsilon}}{\sqrt{n} s_x}$
- σ_{ε} : standard deviation of the errors. Less noise, smaller SE

Standard error of a slope

- $SE(\hat{\beta}) \approx \frac{\sigma_{\varepsilon}}{\sqrt{n} s_x}$
- σ_{ε} : standard deviation of the errors. Less noise, smaller SE
- n : more data, smaller SE

Standard error of a slope

- $SE(\hat{\beta}) \approx \frac{\sigma_{\varepsilon}}{\sqrt{n} s_x}$
- σ_{ε} : standard deviation of the errors. Less noise, smaller SE
- n : more data, smaller SE
- s_x : standard deviation of x . More spread in x , smaller SE

Standard error of a slope

- $SE(\hat{\beta}) \approx \frac{\sigma_{\varepsilon}}{\sqrt{n} s_x}$
- σ_{ε} : standard deviation of the errors. Less noise, smaller SE
- n : more data, smaller SE
- s_x : standard deviation of x . More spread in x , smaller SE
- Then as last week: $t = \hat{\beta} / \widehat{SE}$, 95% CI = $\hat{\beta} \pm 1.96 \times \widehat{SE}$

Assumptions

Assumptions, in order of importance

Assumptions, in order of importance

- Validity: the data map onto the research question

Assumptions, in order of importance

- Validity: the data map onto the research question
- Representativeness: the sample reflects the population

Assumptions, in order of importance

- Validity: the data map onto the research question
- Representativeness: the sample reflects the population
- Additivity and linearity: the model form fits

Assumptions, in order of importance

- Validity: the data map onto the research question
- Representativeness: the sample reflects the population
- Additivity and linearity: the model form fits
- Independence of errors

Assumptions, in order of importance

- Validity: the data map onto the research question
- Representativeness: the sample reflects the population
- Additivity and linearity: the model form fits
- Independence of errors
- Equal variance of errors (homoskedasticity)

Assumptions, in order of importance

- Validity: the data map onto the research question
- Representativeness: the sample reflects the population
- Additivity and linearity: the model form fits
- Independence of errors
- Equal variance of errors (homoskedasticity)
- Normality of errors

The order is the point

The order is the point

- From Gelman, Hill & Vehtari, ch. 11

The order is the point

- From Gelman, Hill & Vehtari, ch. 11
- Most courses spend most time on the last three

The order is the point

- From Gelman, Hill & Vehtari, ch. 11
- Most courses spend most time on the last three
- The first three decide whether $\hat{\beta}$ means what we say

The order is the point

- From Gelman, Hill & Vehtari, ch. 11
- Most courses spend most time on the last three
- The first three decide whether $\hat{\beta}$ means what we say
- The last three mostly affect standard errors

The order is the point

- From Gelman, Hill & Vehtari, ch. 11
- Most courses spend most time on the last three
- The first three decide whether $\hat{\beta}$ means what we say
- The last three mostly affect standard errors
- No software output checks the first two

- Build a data-generating process with a known slope, 0.5

- Build a data-generating process with a known slope, 0.5
- Break one assumption

Simulate, break, observe

- Build a data-generating process with a known slope, 0.5
- Break one assumption
- Fit the regression on 4,000 simulated datasets

Simulate, break, observe

- Build a data-generating process with a known slope, 0.5
- Break one assumption
- Fit the regression on 4,000 simulated datasets
- Bias: is $\hat{\beta}$ centered on 0.5?

- Build a data-generating process with a known slope, 0.5
- Break one assumption
- Fit the regression on 4,000 simulated datasets
- Bias: is $\hat{\beta}$ centered on 0.5?
- Coverage: do 95% intervals contain 0.5 in 95% of datasets?

Four scenarios, $n = 200$

Four scenarios, $n = 200$

- Assumptions hold

Four scenarios, $n = 200$

- Assumptions hold
- Skewed errors: not normal


Four scenarios, $n = 200$

- Assumptions hold
- Skewed errors: not normal
- Unequal variance: errors larger where x is far from its mean

Four scenarios, $n = 200$

- Assumptions hold
- Skewed errors: not normal
- Unequal variance: errors larger where x is far from its mean
- Clustered data: 20 groups of 10, x and part of the error shared within group

Coverage of 95% intervals, 4,000 datasets per scenario



What happened

What happened

- $\hat{\beta}$ centered on 0.5 in all four

What happened

- $\hat{\beta}$ centered on 0.5 in all four
- Skewed errors: no harm at this sample size

What happened

- $\hat{\beta}$ centered on 0.5 in all four
- Skewed errors: no harm at this sample size
- Unequal variance: 75% coverage, robust standard errors restore 94%

What happened

- $\hat{\beta}$ centered on 0.5 in all four
- Skewed errors: no harm at this sample size
- Unequal variance: 75% coverage, robust standard errors restore 94%
- Clustering: 58% coverage, cluster-robust standard errors get to 92%

What happened

- $\hat{\beta}$ centered on 0.5 in all four
- Skewed errors: no harm at this sample size
- Unequal variance: 75% coverage, robust standard errors restore 94%
- Clustering: 58% coverage, cluster-robust standard errors get to 92%
- 20 clusters is few: the correction itself is approximate

- Estimate the SE without assuming equal variance or independence

Robust standard errors

- Estimate the SE without assuming equal variance or independence
- Fix the interval, leave $\hat{\beta}$ untouched

Robust standard errors

- Estimate the SE without assuming equal variance or independence
- Fix the interval, leave $\hat{\beta}$ untouched
- King & Roberts (2015): a large gap between classical and robust SEs is a symptom


Robust standard errors

- Estimate the SE without assuming equal variance or independence
- Fix the interval, leave $\hat{\beta}$ untouched
- King & Roberts (2015): a large gap between classical and robust SEs is a symptom
- Of what: a missing variable, a wrong functional form, ignored grouping

Robust standard errors

- Estimate the SE without assuming equal variance or independence
- Fix the interval, leave $\hat{\beta}$ untouched
- King & Roberts (2015): a large gap between classical and robust SEs is a symptom
- Of what: a missing variable, a wrong functional form, ignored grouping
- Lab 2: clustered and bootstrap standard errors on the ESS

One curve, two samples



When linearity fails

When linearity fails

- OLS still returns the best straight-line approximation

When linearity fails

- OLS still returns the best straight-line approximation
- Which line depends on where the sample's x values sit

When linearity fails

- OLS still returns the best straight-line approximation
- Which line depends on where the sample's x values sit
- Here: slope 1.04 in sample A, -0.44 in sample B, same relationship

When linearity fails

- OLS still returns the best straight-line approximation
- Which line depends on where the sample's x values sit
- Here: slope 1.04 in sample A, -0.44 in sample B, same relationship
- No standard error correction helps

When linearity fails

- OLS still returns the best straight-line approximation
- Which line depends on where the sample's x values sit
- Here: slope 1.04 in sample A, -0.44 in sample B, same relationship
- No standard error correction helps
- Plot the data before trusting a slope

Which assumptions for which purpose

Which assumptions for which purpose

- Description: representativeness, validity, a model form that fits

Which assumptions for which purpose

- Description: representativeness, validity, a model form that fits
- Prediction: fit on new data. Little else matters

Which assumptions for which purpose

- Description: representativeness, validity, a model form that fits
- Prediction: fit on new data. Little else matters
- Causal inference: no confounding, on top of everything description needs

Which assumptions for which purpose

- Description: representativeness, validity, a model form that fits
- Prediction: fit on new data. Little else matters
- Causal inference: no confounding, on top of everything description needs
- Reporting uncertainty in any of them: independence and variance assumptions

Interactions


- $y = \alpha + \beta_1 x + \beta_2 z + \beta_3 xz + \varepsilon$

- $y = \alpha + \beta_1 x + \beta_2 z + \beta_3 xz + \varepsilon$
- The slope of x is no longer one number: $\beta_1 + \beta_3 z$

- $y = \alpha + \beta_1 x + \beta_2 z + \beta_3 xz + \varepsilon$
- The slope of x is no longer one number: $\beta_1 + \beta_3 z$
- z : the moderator

- $y = \alpha + \beta_1 x + \beta_2 z + \beta_3 xz + \varepsilon$
- The slope of x is no longer one number: $\beta_1 + \beta_3 z$
- z : the moderator
- β_1 alone: the slope of x when $z = 0$, nothing more

Slope of x across z , with 95% interval



Rules for interactions

Rules for interactions

- Include x , z and xz : all constitutive terms

Rules for interactions

- Include x , z and xz : all constitutive terms
- Do not read β_1 as an average effect

Rules for interactions

- Include x , z and xz : all constitutive terms
- Do not read β_1 as an average effect
- Plot the slope across z with its interval

Rules for interactions

- Include x , z and xz : all constitutive terms
- Do not read β_1 as an average effect
- Plot the slope across z with its interval
- Brambor, Clark & Golder (2006)

Rules for interactions

- Include x , z and xz : all constitutive terms
- Do not read β_1 as an average effect
- Plot the slope across z with its interval
- Brambor, Clark & Golder (2006)
- Check the interaction is linear and that data exist across z (Hainmueller, Mummolo & Xu 2019)

What to take from this

- Read coefficients as comparisons first

- Read coefficients as comparisons first
- Ask what variation a coefficient is using

- Read coefficients as comparisons first
- Ask what variation a coefficient is using
- Worry about validity and representativeness before normality

- Read coefficients as comparisons first
- Ask what variation a coefficient is using
- Worry about validity and representativeness before normality
- When in doubt about a method, simulate it

Before next week

Before next week

- Lab 2 on Monday: regression in R

Before next week

- Lab 2 on Monday: regression in R
- Week 3, with Sebastian: estimands, DAGs, model specification

Before next week

- Lab 2 on Monday: regression in R
- Week 3, with Sebastian: estimands, DAGs, model specification
- Lundberg, Johnson & Stewart (2021)

Before next week

- Lab 2 on Monday: regression in R
- Week 3, with Sebastian: estimands, DAGs, model specification
- Lundberg, Johnson & Stewart (2021)
- Knox, Lowe & Mummolo (2020)

Before next week

- Lab 2 on Monday: regression in R
- Week 3, with Sebastian: estimands, DAGs, model specification
- Lundberg, Johnson & Stewart (2021)
- Knox, Lowe & Mummolo (2020)
- Bronner (2020)

Questions