

Intermediate Quantitative Methods

The logic of statistical inference

Kevin Munger

Week 1

European University Institute, SPS

Today

- From sample to population

Today

- From sample to population
- Standard errors and intervals

Today

- From sample to population
- Standard errors and intervals
- Hypothesis tests and p-values

Today

- From sample to population
- Standard errors and intervals
- Hypothesis tests and p-values
- Power, and what low power does

From sample to population

The problem

The problem

- We care about a population

The problem

- We care about a population
- We observe a sample

The problem

- We care about a population
- We observe a sample
- How far can one sample take us?

Estimand, estimator, estimate

- Estimand: the quantity we want, e.g. the population mean μ

Estimand, estimator, estimate

- Estimand: the quantity we want, e.g. the population mean μ
- Estimator: the rule we apply, e.g. the sample mean $\bar{x} = \frac{1}{n} \sum_i x_i$

Estimand, estimator, estimate

- Estimand: the quantity we want, e.g. the population mean μ
- Estimator: the rule we apply, e.g. the sample mean $\bar{x} = \frac{1}{n} \sum_i x_i$
- Estimate: the number the rule gives in our sample, 180.05 cm


Estimand, estimator, estimate

- Estimand: the quantity we want, e.g. the population mean μ
- Estimator: the rule we apply, e.g. the sample mean $\bar{x} = \frac{1}{n} \sum_i x_i$
- Estimate: the number the rule gives in our sample, 180.05 cm
- Notation: Greek letters (μ, σ) for the population, Latin letters ($\bar{x}, s$) for the sample


Estimand, estimator, estimate

- Estimand: the quantity we want, e.g. the population mean μ
- Estimator: the rule we apply, e.g. the sample mean $\bar{x} = \frac{1}{n} \sum_i x_i$
- Estimate: the number the rule gives in our sample, 180.05 cm
- Notation: Greek letters (μ, σ) for the population, Latin letters ($\bar{x}, s$) for the sample
- Defining the estimand is often the hard part: week 3

Population: 700,000 simulated heights



One random sample, $n = 859$



Three distributions

Three distributions

- Population: all units we care about, mean μ , standard deviation σ

Three distributions

- Population: all units we care about, mean μ , standard deviation σ
- Sample: the n units we observe, mean $\bar{x}$, standard deviation s


Three distributions

- Population: all units we care about, mean μ , standard deviation σ
- Sample: the n units we observe, mean $\bar{x}$, standard deviation s
- Sampling distribution: the distribution of $\bar{x}$ across repeated samples

Three distributions

- Population: all units we care about, mean μ , standard deviation σ
- Sample: the n units we observe, mean $\bar{x}$, standard deviation s
- Sampling distribution: the distribution of $\bar{x}$ across repeated samples
- Never observed, only derived from theory or simulated

Means of 1,000 samples, $n = 859$



The sampling distribution

The sampling distribution

- Centered on μ : the sample mean is unbiased

The sampling distribution

- Centered on μ : the sample mean is unbiased
- Spread: the standard error (SE)

The sampling distribution


- Centered on μ : the sample mean is unbiased
- Spread: the standard error (SE)
- Shape: approximately normal

The sampling distribution


- Centered on μ : the sample mean is unbiased
- Spread: the standard error (SE)
- Shape: approximately normal
- Here: $\sigma = 7.4$, $n = 859$, so $SE \approx 0.25$ cm

Skewed population, means of 5,000 samples


Population




Means, $n = 2$



Means, $n = 10$



Means, $n = 50$



Central limit theorem

Central limit theorem

- Means of independent draws approach a normal distribution as n grows

Central limit theorem

- Means of independent draws approach a normal distribution as n grows
- Holds for almost any population with finite variance

Central limit theorem

- Means of independent draws approach a normal distribution as n grows
- Holds for almost any population with finite variance
- “ $n \geq 30$ ” is a rule of thumb, not a law

Central limit theorem

- Means of independent draws approach a normal distribution as n grows
- Holds for almost any population with finite variance
- “ $n \geq 30$ ” is a rule of thumb, not a law
- Breaks down with heavy tails, small samples, dependent observations

Central limit theorem

- Means of independent draws approach a normal distribution as n grows
- Holds for almost any population with finite variance
- “ $n \geq 30$ ” is a rule of thumb, not a law
- Breaks down with heavy tails, small samples, dependent observations
- Why normal-based tests work far beyond normal data

Standard errors and intervals

Two kinds of spread

Two kinds of spread

- Standard deviation: spread of individual observations around their mean

Two kinds of spread

- Standard deviation: spread of individual observations around their mean
- $s = \sqrt{\frac{1}{n-1} \sum_i (x_i - \bar{x})^2}$, estimates σ

Two kinds of spread

- Standard deviation: spread of individual observations around their mean
- $s = \sqrt{\frac{1}{n-1} \sum_i (x_i - \bar{x})^2}$, estimates σ
- Standard error: spread of an estimate across repeated samples

Two kinds of spread

- Standard deviation: spread of individual observations around their mean
- $s = \sqrt{\frac{1}{n-1} \sum_i (x_i - \bar{x})^2}$, estimates σ
- Standard error: spread of an estimate across repeated samples
- Easy to confuse, answer different questions


- $SE = \sigma / \sqrt{n}$

- $SE = \sigma / \sqrt{n}$
- More variation in the population, larger SE

- $SE = \sigma / \sqrt{n}$
- More variation in the population, larger SE
- Larger sample, smaller SE, but only by $\sqrt{n}$: 4 times the data halves the SE

- $SE = \sigma / \sqrt{n}$
- More variation in the population, larger SE
- Larger sample, smaller SE, but only by $\sqrt{n}$: 4 times the data halves the SE
- σ unknown, so we plug in s : $\widehat{SE} = s / \sqrt{n}$

Standard error by sample size, $\sigma = 7.4$



The normal distribution

The normal distribution

- Symmetric and bell-shaped, defined by its mean and standard deviation

The normal distribution

- Symmetric and bell-shaped, defined by its mean and standard deviation
- z-score: distance from the mean in standard deviations, $z = (x - \mu)/\sigma$

The normal distribution

- Symmetric and bell-shaped, defined by its mean and standard deviation
- z-score: distance from the mean in standard deviations, $z = (x - \mu)/\sigma$
- Standard normal $N(0, 1)$: the distribution of z

The normal distribution

- Symmetric and bell-shaped, defined by its mean and standard deviation
- z-score: distance from the mean in standard deviations, $z = (x - \mu)/\sigma$
- Standard normal $N(0, 1)$: the distribution of z
- 68% of the mass within ± 1 , 95% within ± 1.96 , 99% within ± 2.58

- 95% CI: $\bar{x} \pm 1.96 \times \widehat{SE}$

Confidence intervals

- 95% CI: $\bar{x} \pm 1.96 \times \widehat{SE}$
- 1.96 comes from the standard normal

Confidence intervals

- 95% CI: $\bar{x} \pm 1.96 \times \widehat{SE}$
- 1.96 comes from the standard normal
- t distribution: like the standard normal, with fatter tails

Confidence intervals

- 95% CI: $\bar{x} \pm 1.96 \times \widehat{SE}$
- 1.96 comes from the standard normal
- t distribution: like the standard normal, with fatter tails
- Fatter tails account for estimating σ with s

Confidence intervals

- 95% CI: $\bar{x} \pm 1.96 \times \widehat{SE}$
- 1.96 comes from the standard normal
- t distribution: like the standard normal, with fatter tails
- Fatter tails account for estimating σ with s
- Depends on degrees of freedom ($n - 1$): 2.09 at $n = 20$, 1.98 at $n = 100$, then ≈ 1.96

Example: the sample of 859

Example: the sample of 859

- $\bar{x} = 180.05$ cm, $s = 7.52$

Example: the sample of 859

- $\bar{x} = 180.05$ cm, $s = 7.52$
- $\widehat{SE} = 7.52/\sqrt{859} = 0.26$


Example: the sample of 859

- $\bar{x} = 180.05 \text{ cm}, s = 7.52$
- $\widehat{SE} = 7.52/\sqrt{859} = 0.26$
- 95% CI: $180.05 \pm 1.96 \times 0.26 = [179.55, 180.56]$

Example: the sample of 859

- $\bar{x} = 180.05$ cm, $s = 7.52$
- $\widehat{SE} = 7.52/\sqrt{859} = 0.26$
- 95% CI: $180.05 \pm 1.96 \times 0.26 = [179.55, 180.56]$
- $\mu = 180.15$: inside, this time

100 samples of 100, 100 95% intervals



What 95% refers to

What 95% refers to

- The procedure, not the interval

What 95% refers to

- The procedure, not the interval
- 95% of intervals from repeated samples contain μ

What 95% refers to

- The procedure, not the interval
- 95% of intervals from repeated samples contain μ
- Any single interval either does or does not

What 95% refers to

- The procedure, not the interval
- 95% of intervals from repeated samples contain μ
- Any single interval either does or does not
- Not a 95% probability that μ is in this interval

Hypothesis tests and p-values

- Null hypothesis H_0 : the conservative claim, usually no difference

The logic

- Null hypothesis H_0 : the conservative claim, usually no difference
- Alternative H_A : the research hypothesis, some difference

- Null hypothesis H_0 : the conservative claim, usually no difference
- Alternative H_A : the research hypothesis, some difference
- Test statistic: distance of the estimate from the null, in SEs, $t = (\hat{\theta} - \theta_0) / \widehat{SE}$

- Null hypothesis H_0 : the conservative claim, usually no difference
- Alternative H_A : the research hypothesis, some difference
- Test statistic: distance of the estimate from the null, in SEs, $t = (\hat{\theta} - \theta_0) / \widehat{SE}$
- p-value: how often a test statistic this extreme would arise if H_0 were true

The logic

- Null hypothesis H_0 : the conservative claim, usually no difference
- Alternative H_A : the research hypothesis, some difference
- Test statistic: distance of the estimate from the null, in SEs, $t = (\hat{\theta} - \theta_0) / \widehat{SE}$
- p-value: how often a test statistic this extreme would arise if H_0 were true
- Reject H_0 if $p < \alpha$, by convention $\alpha = 0.05$

Example: age in northern and southern Italy

Example: age in northern and southern Italy

- ESS 2016, 2,626 respondents

Example: age in northern and southern Italy

- ESS 2016, 2,626 respondents
- $H_0: \mu_{\text{North}} - \mu_{\text{South}} = 0$

Example: age in northern and southern Italy

- ESS 2016, 2,626 respondents
- $H_0: \mu_{\text{North}} - \mu_{\text{South}} = 0$
- Estimated difference 2.44 years, $\widehat{SE} = 0.79$


Example: age in northern and southern Italy

- ESS 2016, 2,626 respondents
- $H_0: \mu_{\text{North}} - \mu_{\text{South}} = 0$
- Estimated difference 2.44 years, $\widehat{SE} = 0.79$
- $t = 2.44/0.79 \approx 3.08$

Example: age in northern and southern Italy

- ESS 2016, 2,626 respondents
- $H_0: \mu_{\text{North}} - \mu_{\text{South}} = 0$
- Estimated difference 2.44 years, $\widehat{SE} = 0.79$
- $t = 2.44/0.79 \approx 3.08$
- $p = 0.002$: about 2 in 1,000 samples under H_0

Null distribution, $t = \pm 3.08$



What a p-value is

What a p-value is

- $P(\text{data at least this extreme} \mid H_0 \text{ true})$

What a p-value is

- $P(\text{data at least this extreme} \mid H_0 \text{ true})$
- Not $P(H_0 \text{ true})$

What a p-value is

- $P(\text{data at least this extreme} \mid H_0 \text{ true})$
- Not $P(H_0 \text{ true})$
- Not the size of the effect


What a p-value is

- $P(\text{data at least this extreme} \mid H_0 \text{ true})$
- Not $P(H_0 \text{ true})$
- Not the size of the effect
- Not the probability of replication

What a p-value is

- $P(\text{data at least this extreme} \mid H_0 \text{ true})$
- Not $P(H_0 \text{ true})$
- Not the size of the effect
- Not the probability of replication
- Under H_0 , every p-value between 0 and 1 is equally likely

p-values from 10,000 simulated studies, $n = 100$ per group



Two kinds of error

Two kinds of error

- Type I: reject a true H_0 , a false positive

Two kinds of error

- Type I: reject a true H_0 , a false positive
- Type II: fail to reject a false H_0 , a false negative

Two kinds of error

- Type I: reject a true H_0 , a false positive
- Type II: fail to reject a false H_0 , a false negative
- α : the Type I error rate we accept, 0.05 by convention

Two kinds of error

- Type I: reject a true H_0 , a false positive
- Type II: fail to reject a false H_0 , a false negative
- α : the Type I error rate we accept, 0.05 by convention
- β : the Type II error rate

Two kinds of error

- Type I: reject a true H_0 , a false positive
- Type II: fail to reject a false H_0 , a false negative
- α : the Type I error rate we accept, 0.05 by convention
- β : the Type II error rate
- Power = $1 - \beta$: probability of detecting an effect that is there

Power

- Cohen's d : difference in means divided by the standard deviation

Effect size and power

- Cohen's d : difference in means divided by the standard deviation
- Conventions: 0.2 small, 0.5 medium, 0.8 large

Effect size and power

- Cohen's d : difference in means divided by the standard deviation
- Conventions: 0.2 small, 0.5 medium, 0.8 large
- Power depends on d , n and α


Effect size and power

- Cohen's d : difference in means divided by the standard deviation
- Conventions: 0.2 small, 0.5 medium, 0.8 large
- Power depends on d , n and α
- Fix three, solve for the fourth: a power analysis

Effect size and power

- Cohen's d : difference in means divided by the standard deviation
- Conventions: 0.2 small, 0.5 medium, 0.8 large
- Power depends on d , n and α
- Fix three, solve for the fourth: a power analysis
- Convention: aim for 80% power

Power by sample size, two groups, $\alpha = 0.05$



How much power does political science have?

How much power does political science have?

- 16,000 tests from about 2,000 articles

How much power does political science have?

- 16,000 tests from about 2,000 articles
- Power to detect the typical effect size in each literature

How much power does political science have?

- 16,000 tests from about 2,000 articles
- Power to detect the typical effect size in each literature
- Median power about 10%


How much power does political science have?

- 16,000 tests from about 2,000 articles
- Power to detect the typical effect size in each literature
- Median power about 10%
- About 1 in 10 tests reach 80%

How much power does political science have?

- 16,000 tests from about 2,000 articles
- Power to detect the typical effect size in each literature
- Median power about 10%
- About 1 in 10 tests reach 80%
- Arel-Bundock et al. (2026)

5,000 studies, true effect 0.1, $n = 100$ per group



What low power does

What low power does

- Most real effects are missed

What low power does

- Most real effects are missed
- Type M (magnitude) error: significant estimates exaggerate the true effect

What low power does

- Most real effects are missed
- Type M (magnitude) error: significant estimates exaggerate the true effect
- Type S (sign) error: significant estimates with the wrong sign

What low power does

- Most real effects are missed
- Type M (magnitude) error: significant estimates exaggerate the true effect
- Type S (sign) error: significant estimates with the wrong sign
- Here: power 11%, significant estimates about 3 times the truth, 3% wrong sign

What low power does

- Most real effects are missed
- Type M (magnitude) error: significant estimates exaggerate the true effect
- Type S (sign) error: significant estimates with the wrong sign
- Here: power 11%, significant estimates about 3 times the truth, 3% wrong sign
- Gelman & Carlin (2014)

Of the significant results, how many are true?

Of the significant results, how many are true?

- False discovery rate: share of significant results that are false

Of the significant results, how many are true?

- False discovery rate: share of significant results that are false
- Depends on α


Of the significant results, how many are true?

- False discovery rate: share of significant results that are false
- Depends on α
- Depends on power

Of the significant results, how many are true?

- False discovery rate: share of significant results that are false
- Depends on α
- Depends on power
- Depends on the share of tested hypotheses that are true

1,000 hypotheses: 30% true, 80% power, $\alpha = 0.05$



The same 1,000 hypotheses

The same 1,000 hypotheses

- 80% power: 240 true positives, 35 false positives

The same 1,000 hypotheses

- 80% power: 240 true positives, 35 false positives
- 10% power: 30 true positives, 35 false positives

The same 1,000 hypotheses

- 80% power: 240 true positives, 35 false positives
- 10% power: 30 true positives, 35 false positives
- More false positives than true ones

- Many defensible analyses, one reported

Forking paths

- Many defensible analyses, one reported
- 20 outcomes, no true effects

- Many defensible analyses, one reported
- 20 outcomes, no true effects
- $P(\text{at least one } p < 0.05) = 1 - 0.95^{20} \approx 0.64$

- Many defensible analyses, one reported
- 20 outcomes, no true effects
- $P(\text{at least one } p < 0.05) = 1 - 0.95^{20} \approx 0.64$
- Gelman & Loken (2014), bonus exercise in Lab 1

What to take from this

Significance is not importance

Significance is not importance

- Large samples make trivial effects significant

Significance is not importance

- Large samples make trivial effects significant
- A large p-value is not evidence of no effect

Significance is not importance

- Large samples make trivial effects significant
- A large p-value is not evidence of no effect
- Report estimates and intervals, not stars

- Think about power before collecting data

- Think about power before collecting data
- Simulate the design

- Think about power before collecting data
- Simulate the design
- Treat 0.05 as a convention

- Think about power before collecting data
- Simulate the design
- Treat 0.05 as a convention
- Distrust large effects from small studies

Before next week

Before next week

- Install R, Positron or RStudio, Quarto

Before next week

- Install R, Positron or RStudio, Quarto
- First lab on Monday

Before next week

- Install R, Positron or RStudio, Quarto
- First lab on Monday
- Gelman, Hill & Vehtari, *Regression and Other Stories*, ch. 6–8, 10

Before next week

- Install R, Positron or RStudio, Quarto
- First lab on Monday
- Gelman, Hill & Vehtari, *Regression and Other Stories*, ch. 6–8, 10
- Munger, Guess & Hargittai (2021)

Questions